

Artificial Intelligence and Machine Learning for Connected Industries

Class 12 - TensorFlow

Killian Cressant

Introduction

TensorFlow lower-level API

- 95% of the use cases you will encounter
 - Keras is enough
- It will be useful when we need extra control to write custom
 - Loss functions, metrics, layers, models, initializers, regularizers, weight constraints
 - To fully control the training loop itself
 - Apply special transformations or constraints to the gradients
 - Use multiple optimizers for different parts of the network

TensorFlow

- TensorFlow is a powerful library for numerical computation
 - Particularly well suited and fine-tuned for large-scale Machine Learning
- Developed by the Google Brain team
- It powers many of Google's large-scale services
 - Google Cloud Speech, Google Photos, and Google Search


Functionalities

- Its core is very similar to NumPy but with GPU support
- It supports distributed computing across multiple devices and servers
- It includes a kind of just-in-time (JIT) compiler
 - Allows it to optimize computations for speed and memory usage
- Computation graphs can be exported to a portable format
- It provides some excellent optimizers
 - RMSProp and Nadam
 - Can easily minimize all sorts of loss functions

Functionalities

- Keras
 - tf.keras
- Loading and processing
 - tf.data and tf.io
- Image processing
 - tf.image
- Signal processing
 - tf.signal

TensorFlow's Python API




source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

TensorFlow operations

- At the lowest level
 - Each TensorFlow operation (op for short) is implemented
 - Using highly efficient C++ code
- Many operations have multiple implementations called kernels
 - Each kernel is dedicated to a specific device type
 - CPUs
 - GPUs
 - TPUs —> Tensor Processing Units

TensorFlow architecture



source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Multi-platform

- TensorFlow runs on multiple operating systems
 - Windows, Linux, and macOS
 - Mobile devices —> using TensorFlow Lite
 - iOS and Android
- APIs for other languages are also available
 - C++, Java, and Swift APIs
 - JavaScript implementation called TensorFlow.js
 - Makes it possible to run your models directly in your browser

An ecosystem of libraries

- TensorBoard for visualization
- TensorFlow Extended (TFX)
 - A set of libraries built by Google
 - Productionize TensorFlow projects
 - It includes tools for
 - Data validation and preprocessing
 - Model analysis

An ecosystem of libraries

- Google's TensorFlow Hub
 - Easily download and reuse pre-trained neural networks
 - Model garden
 - We can get many neural network architectures
 - Some of them pre-trained

<https://paperswithcode.com/>

Using TensorFlow like NumPy

TensorFlow like NumPy

- TensorFlow's API revolves around tensors
 - Which flow from operation to operation
 - A tensor is very similar to a NumPy ndarray
 - It is usually a multidimensional array
 - It can also hold a scalar
 - A simple value, such as 27

Tensor example

```
import tensorflow as tf  
  
t = tf.constant([[1., 2., 3.], [4., 5., 6.]]) # matrix  
t
```

```
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=  
array([[1., 2., 3.],  
       [4., 5., 6.]]), dtype=float32>
```

```
>>> t.shape  
TensorShape([2, 3])  
>>> t.dtype  
tf.float32
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Indexing

```
>>> t[:, 1:]  
<tf.Tensor: shape=(2, 2), dtype=float32, numpy=  
array([[2., 3.],  
       [5., 6.]]), dtype=float32>
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Simple operations

```
>>> t + 10  
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=  
array([[11., 12., 13.],  
       [14., 15., 16.]], dtype=float32)>  
>>> tf.square(t)  
<tf.Tensor: shape=(2, 3), dtype=float32, numpy=  
array([[ 1.,   4.,   9.],  
       [16., 25., 36.]], dtype=float32)>  
>>> t @ tf.transpose(t)  
<tf.Tensor: shape=(2, 2), dtype=float32, numpy=  
array([[14., 32.],  
       [32., 77.]], dtype=float32)>
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Simple operations

- Most operations are the same in NumPy
- Some have different names
 - `tf.reduce_mean()` —> `np.mean()`
 - `tf.reduce_sum()` —> `np.sum()`
 - `tf.reduce_max()` —> `np.max()`
 - `tf.math.log()` —> `np.log()`

Tensors and NumPy

- We can create a tensor from a NumPy array, and vice versa
- We can even apply TensorFlow operations to NumPy arrays
 - And NumPy operations to tensors
- Precision by default
 - NumPy —> uses 64-bit
 - TensorFlow —> 32-bit
 - It is generally more than enough for neural networks
 - It runs faster and uses less RAM

```
set dtype = tf.float32
```

Type conversions

- Type conversions can significantly hurt performance
- They can easily go unnoticed when they are done automatically
- TensorFlow does not perform any type conversions automatically
 - It just raises an exception
 - When we try to execute an operation on tensors with incompatible types

Variables

- The `tf.Tensor` values we've seen so far are immutable
 - This means that we cannot use regular tensors
 - To implement weights in a neural network
 - Since they need to be tweaked by back-propagation
 - Other parameters may also need to change over time

```
>>> v = tf.Variable([[1., 2., 3.], [4., 5., 6.]])  
>>> v  
<tf.Variable 'Variable:0' shape=(2, 3) dtype=float32, numpy=  
array([[1., 2., 3.],  
       [4., 5., 6.]], dtype=float32)>
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Variables

```
v.assign(2 * v)          # v now equals [[2., 4., 6.], [8., 10., 12.]]  
v[0, 1].assign(42)       # v now equals [[2., 42., 6.], [8., 10., 12.]]  
v[:, 2].assign([0., 1.]) # v now equals [[2., 42., 0.], [8., 10., 1.]]  
v.scatter_nd_update(  
    indices=[[0, 0], [1, 2]], updates=[100., 200.])
```

- Direct assignment will not work

```
>>> v[1] = [7., 8., 9.]  
[...] TypeError: 'ResourceVariable' object does not support item assignment
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Customizing Models and Training Algorithms

Supposition

- Suppose we want to train a regression model
 - Our training set is a bit noisy
- We try to clean up our dataset
 - Eliminating the outliers
 - It is not so efficient
- Which loss function we are going to use?
 - MSE
 - Absolute error


Custom Loss Functions

- Huber loss function
 - Imagine it is not implemented in Keras
 - `tf.keras.losses.Huber` class

```
def huber_fn(y_true, y_pred):  
    error = y_true - y_pred  
    is_small_error = tf.abs(error) < 1  
    squared_loss = tf.square(error) / 2  
    linear_loss = tf.abs(error) - 0.5  
    return tf.where(is_small_error, squared_loss, linear_loss)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Custom Huber loss



source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Using our custom loss function

- Now we can use this Huber loss function
 - When we compile the Keras model
 - Train your model as usual

```
model.compile(loss=huber_fn, optimizer="nadam")
model.fit(X_train, y_train, [...])
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Saving and loading our model

- When we have a custom function
 - Loading the model
 - We need to provide a dictionary
 - Maps the function name to the actual function

```
model = tf.keras.models.load_model("my_model_with_a_custom_loss",
                                    custom_objects={"huber_fn": huber_fn})
```

- If you decorate the `huber_fn()` function with
 - `@keras.utils.register_keras_serializable()`

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

The threshold as a parameter

```
def create_huber(threshold=1.0):
    def huber_fn(y_true, y_pred):
        error = y_true - y_pred
        is_small_error = tf.abs(error) < threshold
        squared_loss = tf.square(error) / 2
        linear_loss = threshold * tf.abs(error) - threshold ** 2 / 2
        return tf.where(is_small_error, squared_loss, linear_loss)
    return huber_fn

model.compile(loss=create_huber(2.0), optimizer="nadam")
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Loading our model

```
model = tf.keras.models.load_model(  
    "my_model_with_a_custom_loss_threshold_2",  
    custom_objects={"huber_fn": create_huber(2.0)}  
)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

An alternative way

```
class HuberLoss(tf.keras.losses.Loss):  
    def __init__(self, threshold=1.0, **kwargs):  
        self.threshold = threshold  
        super().__init__(**kwargs)  
  
    def call(self, y_true, y_pred):  
        error = y_true - y_pred  
        is_small_error = tf.abs(error) < self.threshold  
        squared_loss = tf.square(error) / 2  
        linear_loss = self.threshold * tf.abs(error) - self.threshold**2 / 2  
        return tf.where(is_small_error, squared_loss, linear_loss)  
  
    def get_config(self):  
        base_config = super().get_config()  
        return {**base_config, "threshold": self.threshold}
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Custom Activation Functions, Initializers, Regularizers, and Constraints

```
def my_softplus(z):
    return tf.math.log(1.0 + tf.exp(z))

def my_glorot_initializer(shape, dtype=tf.float32):
    stddev = tf.sqrt(2. / (shape[0] + shape[1]))
    return tf.random.normal(shape, stddev=stddev, dtype=dtype)

def my_l1_regularizer(weights):
    return tf.reduce_sum(tf.abs(0.01 * weights))
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Using our custom function

```
layer = tf.keras.layers.Dense(1, activation=my_softplus,  
                             kernel_initializer=my_glorot_initializer,  
                             kernel_regularizer=my_l1_regularizer,  
                             kernel_constraint=my_positive_weights)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Loading and Preprocessing Data with TensorFlow

Transforming items

- Use the map() method

```
dataset = dataset.map(lambda x: x * 2) # x is a batch
for item in dataset:
    print(item)
```

```
tf.Tensor([ 0  2  4  6  8 10 12], shape=(7,), dtype=int32)
tf.Tensor([14 16 18  0  2  4  6], shape=(7,), dtype=int32)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

- We can run using multiple threads to speed things up
 - Setting the num_parallel_calls argument to the number of threads to run

Filtering the dataset

- Use the filter() method

```
dataset = dataset.filter(lambda x: tf.reduce_sum(x) > 50)
for item in dataset:
    print(item)
```

```
tf.Tensor([14 16 18  0  2  4  6], shape=(7,), dtype=int32)
tf.Tensor([ 8 10 12 14 16 18  0], shape=(7,), dtype=int32)
tf.Tensor([ 2  4  6  8 10 12 14], shape=(7,), dtype=int32)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Taking a look at some items

- Using the `take()` method

```
for item in dataset.take(2):  
    print(item)
```

```
tf.Tensor([14 16 18  0  2  4  6], shape=(7,), dtype=int32)  
tf.Tensor([ 8 10 12 14 16 18  0], shape=(7,), dtype=int32)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Shuffling the dataset

```
dataset = tf.data.Dataset.range(10).repeat(2)
dataset = dataset.shuffle(buffer_size=4, seed=42).batch(7)
for item in dataset:
    print(item)

tf.Tensor([3 0 1 6 2 5 7], shape=(7,), dtype=int64)
tf.Tensor([8 4 1 9 4 2 3], shape=(7,), dtype=int64)
tf.Tensor([7 5 0 8 9 6], shape=(6,), dtype=int64)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Shuffling the dataset

- For a large dataset that does not fit in memory
 - This simple shuffling-buffer approach may not be sufficient
 - Since the buffer will be small compared to the dataset
- We can divide the dataset into multiple files
 - Multiple files randomly and read them simultaneously
 - Interleaving their records

```
n_readers = 5  
  
dataset = filepath_dataset.interleave(  
    lambda filepath: tf.data.TextLineDataset(filepath).skip(1),  
    cycle_length=n_readers)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Visualizing some of the dataset

```
for line in dataset.take(5):  
    print(line)
```

```
tf.Tensor(b'4.5909,16.0,[...],33.63,-117.71,2.418', shape=(), dtype=string)  
tf.Tensor(b'2.4792,24.0,[...],34.18,-118.38,2.0', shape=(), dtype=string)  
tf.Tensor(b'4.2708,45.0,[...],37.48,-122.19,2.67', shape=(), dtype=string)  
tf.Tensor(b'2.1856,41.0,[...],32.76,-117.12,1.205', shape=(), dtype=string)  
tf.Tensor(b'4.1812,52.0,[...],33.73,-118.31,3.215', shape=(), dtype=string)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Preprocessing the Data

Scaling

```
x_mean, x_std = [...]
n_inputs = 8

def parse_csv_line(line):
    defs = [0.] * n_inputs + [tf.constant([], dtype=tf.float32)]
    fields = tf.io.decode_csv(line, record_defaults=defs)
    return tf.stack(fields[:-1]), tf.stack(fields[-1:])

def preprocess(line):
    x, y = parse_csv_line(line)
    return (x - x_mean) / x_std, y
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Scaling

```
x_mean, x_std = [...]
n_inputs = 8

def parse_csv_line(line):
    defs = [0.] * n_inputs + [tf.constant([], dtype=tf.float32)]
    fields = tf.io.decode_csv(line, record_defaults=defs)
    return tf.stack(fields[:-1]), tf.stack(fields[-1:])

def preprocess(line):
    x, y = parse_csv_line(line)
    return (x - x_mean) / x_std, y
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Scaling

```
x_mean, x_std = [...]
n_inputs = 8

def parse_csv_line(line):
    defs = [0.] * n_inputs + [tf.constant([], dtype=tf.float32)]
    fields = tf.io.decode_csv(line, record_defaults=defs)
    return tf.stack(fields[:-1]), tf.stack(fields[-1:])

def preprocess(line):
    x, y = parse_csv_line(line)
    return (x - x_mean) / x_std, y
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Scaling

```
x_mean, x_std = [...]
n_inputs = 8

def parse_csv_line(line):
    defs = [0.] * n_inputs + [tf.constant([], dtype=tf.float32)]
    fields = tf.io.decode_csv(line, record_defaults=defs)
    return tf.stack(fields[:-1]), tf.stack(fields[-1:])

def preprocess(line):
    x, y = parse_csv_line(line)
    return (x - x_mean) / x_std, y
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"


Applying to one line

```
preprocess(b'4.2083,44.0,5.3232,0.9171,846.0,2.3370,37.47,-122.2,2.782')
```

```
(<tf.Tensor: shape=(8,), dtype=float32, numpy=
array([ 0.16579159,  1.216324   , -0.05204564, -0.39215982, -0.5277444 ,
       -0.2633488 ,  0.8543046 , -1.3072058 ], dtype=float32)>,
<tf.Tensor: shape=(1,), dtype=float32, numpy=array([2.782], dtype=float32)>
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Loading and preprocessing data from multiple CSV files



source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Using with Keras

```
train_set = csv_reader_dataset(train_filepaths)
valid_set = csv_reader_dataset(valid_filepaths)
test_set = csv_reader_dataset(test_filepaths)
```

```
model = tf.keras.Sequential([...])
model.compile(loss="mse", optimizer="sgd")
model.fit(train_set, validation_data=valid_set, epochs=5)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Processing layers with Keras

Keras Preprocessing Layers


- We can include preprocessing layers directly inside our model
 - It can preprocess all the input data on the fly during training
 - Use the same preprocessing layers in production
- Keras offers many preprocessing layers
 - Can be applied to
 - Numerical features
 - Categorical features
 - Images
 - Text

The Normalization Layer

```
norm_layer = tf.keras.layers.Normalization()
model = tf.keras.models.Sequential([
    norm_layer,
    tf.keras.layers.Dense(1)
])
model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(learning_rate=2e-2))
norm_layer.adapt(x_train) # computes the mean and variance of every feature
model.fit(x_train, y_train, validation_data=(x_valid, y_valid), epochs=5)
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Including preprocessing layers inside a model



Scaling before training

```
norm_layer = tf.keras.layers.Normalization()  
norm_layer.adapt(X_train)  
X_train_scaled = norm_layer(X_train)  
X_valid_scaled = norm_layer(X_valid)
```


```
model = tf.keras.models.Sequential([tf.keras.layers.Dense(1)])  
model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(learning_rate=2e  
model.fit(X_train_scaled, y_train, epochs=5,  
          validation_data=(X_valid_scaled, y_valid))
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Different final model

```
final_model = tf.keras.Sequential([norm_layer, model])
x_new = x_test[:3] # pretend we have a few new instances (unscaled)
y_pred = final_model(x_new) # preprocesses the data and makes predictions
```

Different final model



source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

The Discretization Layer

- Transforms a numerical feature into a categorical feature
 - Mapping value ranges (bins) to categories
- This is sometimes useful for features with
 - Multimodal distributions
 - Have a highly non-linear relationship with the target

Example

```
age = tf.constant([[10.], [93.], [57.], [18.], [37.], [5.]])  
discretize_layer = tf.keras.layers.Discretization(bin_boundaries=[18., 5  
age_categories = discretize_layer(age)  
age_categories  
  
<tf.Tensor: shape=(6, 1), dtype=int64, numpy=array([[0],[2],[2],[1],[1],[0])
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Another example

- We can provide the number of bins
 - Call the layer's adapt() method
 - Let it find the appropriate bin boundaries
 - Based on the value percentiles

```
discretize_layer = tf.keras.layers.Discretization(num_bins=3)
discretize_layer.adapt(age)
age_categories = discretize_layer(age)
age_categories
```

```
<tf.Tensor: shape=(6, 1), dtype=int64, numpy=array([[1],[2],[2],[1],[2],[0]])
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

The CategoryEncoding Layer

- When there are only a few categories
 - Less than a dozen or two
 - The one-hot encoding is often a good option
- Keras provides the CategoryEncoding layer

One-hot-encoding example

```
onehot_layer = tf.keras.layers.CategoryEncoding(num_tokens=3)
onehot_layer(age_categories)
```

```
<tf.Tensor: shape=(6, 3), dtype=float32, numpy=
array([[0., 1., 0.],
       [0., 0., 1.],
       [0., 0., 1.],
       [0., 1., 0.],
       [0., 0., 1.],
       [1., 0., 0.]], dtype=float32)>
```

source: Géron, A. "Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow"

Other layers

- The StringLookup Layer
- The Hashing Layer
- The Embedding Layer
- TextVectorization
- Among others

Artificial Intelligence and Machine Learning for Connected Industries

Class 12 - TensorFlow

Killian Cressant